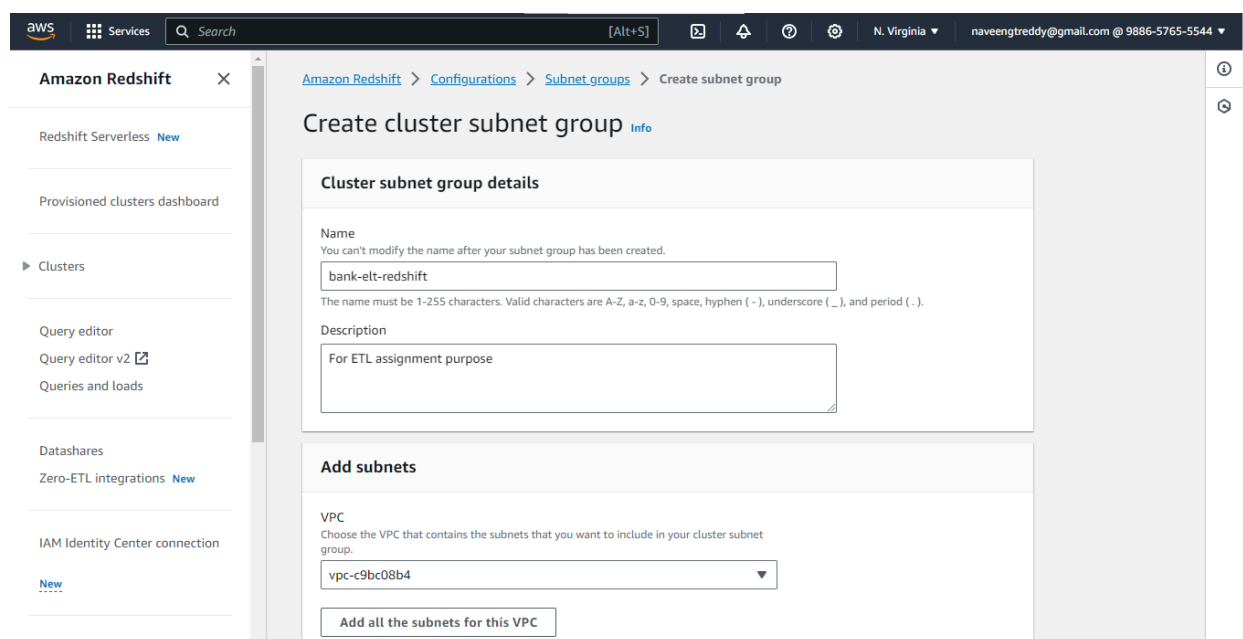


# Creation of a Redshift Cluster

Screenshots of the configuration of the Redshift cluster that you have created:

Create cluster subnet group -




**Amazon Redshift** ×

Redshift Serverless **New**

Provisioned clusters dashboard

► Clusters

Query editor

Query editor v2 

Queries and loads

Datashares

Zero-ETL integrations **New**

IAM Identity Center connection

**New**

Amazon Redshift > Configurations > Subnet groups > Create subnet group

## Create cluster subnet group [Info](#)

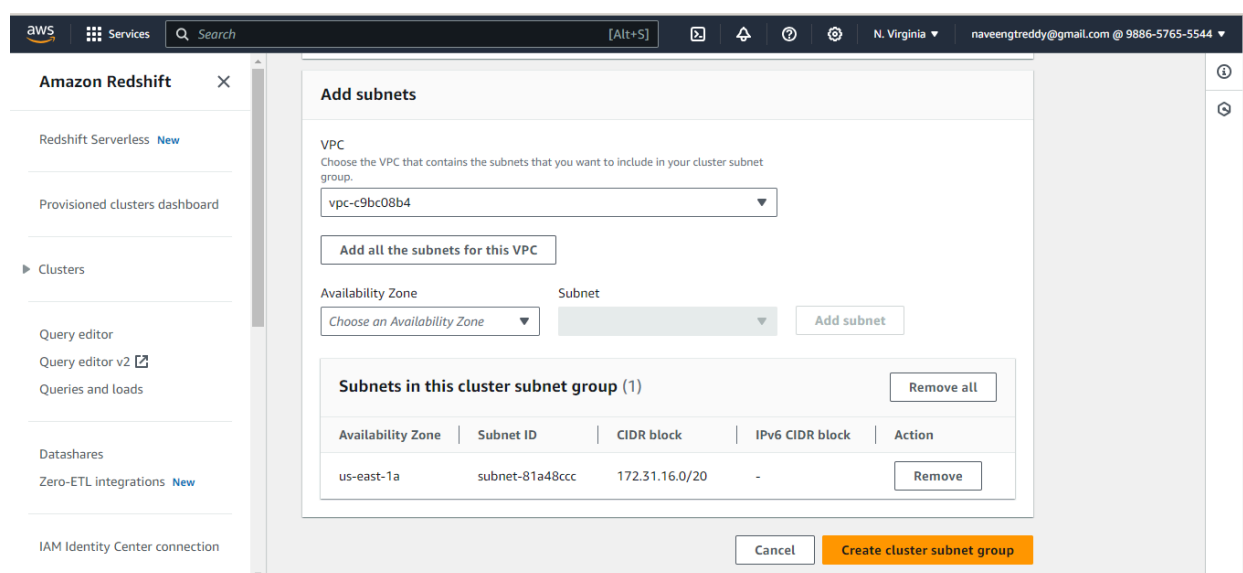
**Cluster subnet group details**

**Name**  
You can't modify the name after your subnet group has been created.  
  
The name must be 1-255 characters. Valid characters are A-Z, a-z, 0-9, space, hyphen ( - ), underscore ( \_ ), and period ( . ).

**Description**

**Add subnets**

**VPC**  
Choose the VPC that contains the subnets that you want to include in your cluster subnet group.




**Amazon Redshift** ×

Redshift Serverless **New**

Provisioned clusters dashboard

► Clusters

Query editor

Query editor v2 

Queries and loads

Datashares

Zero-ETL integrations **New**

IAM Identity Center connection

**Add subnets**

**VPC**  
Choose the VPC that contains the subnets that you want to include in your cluster subnet group.

**Availability Zone**  **Subnet**

**Subnets in this cluster subnet group (1)**

Availability Zone	Subnet ID	CIDR block	IPv6 CIDR block	Action
us-east-1a	subnet-81a48ccc	172.31.16.0/20	-	<input type="button" value="Remove"/>

Cluster subnet group bank-elt-redshift was created successfully

Amazon Redshift > Configurations > Subnet groups

Cluster subnet groups (1) [Info](#)

<input type="checkbox"/>	Name	Status	VPC ID	Description	Tags
<input type="checkbox"/>	bank-elt-redshift 1 Subnets	Complete	vpc-c9bc08b4	For ETL assignment purpose	

## Creating redshift cluster

Node type and number of nodes -

Amazon Redshift > Clusters > Create cluster

### Create cluster [Info](#)

**Cluster configuration**

**Cluster identifier**  
This is the unique key that identifies a cluster.  
  
The identifier must be from 1-63 characters. Valid characters are a-z (lowercase only) and - (hyphen).

**Choose the size of the cluster**  
☒ I'll choose  
☐ Help me choose

**Node type** [Info](#)  
Choose a node type that meets your CPU, RAM, storage capacity, and drive type requirements.

**Number of nodes**  
Enter the number of nodes that you need.  
  
 Range (1-32)

## Configuring database -

aws

Services

Search

[Alt+S]

📧

🔔

🔍

⚙️

N. Virginia ▼

naveentreddy@gmail.com @ 9886-5765-5544 ▼

Amazon Redshift

×

Redshift Serverless New

Provisioned clusters dashboard

Clusters

Reserved nodes

Snapshots

Query editor

Query editor v2

Queries and loads

Cluster permissions

ⓘ

Associated IAM roles (1) Info

Set default ▼

Manage IAM roles ▼

Create, associate, or remove an IAM role. You can associate up to 50 IAM roles. You can also choose an IAM role and set it as the default for this cluster.

🔍 Search for associated IAM role by name, status, or role type

< 1 >

<input type="checkbox"/>	IAM roles <span>🔗</span>	Status	Role type
<input type="checkbox"/>	AmazonRedshift-CommandsAccessRole-20240204T133133	Not applied	Default

aws

Services

Search

[Alt+S]

🔍

🔔

?

⚙️

N. Virginia ▼ naveengtreddy@gmail.com @ 9886-5765-5544 ▼

Amazon Redshift

✕

Redshift Serverless New

Provisioned clusters dashboard

Clusters

- Reserved nodes
- Snapshots

Query editor

Query editor v2 📄

Queries and loads

Dashshares

Zero-ETL integrations New

▼ Database configurations info

Database name

Specify a database name to create an additional database.

redshiftdb

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word. 🔗

Database port

Port number where the database accepts inbound connections. The default port is 5439.

5439

Choose a port number between 1150 and 65535.

Parameter groups

Defines database parameter and query queues for all the databases.

default.redshift-1.0 ▼

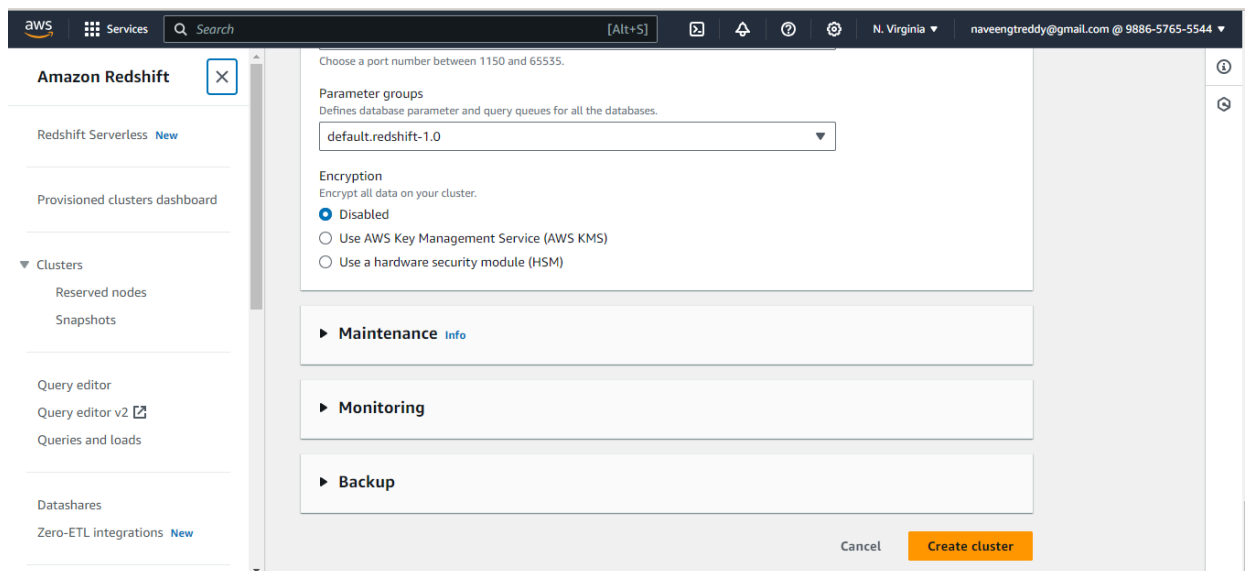
Encryption

Encrypt all data on your cluster.

☒ Disabled

☐ Use AWS Key Management Service (AWS KMS)

☐ Use a hardware security module (HSM)



Choose a port number between 1150 and 65535.

**Parameter groups**  
Defines database parameter and query queues for all the databases.  
default.redshift-1.0

**Encryption**  
Encrypt all data on your cluster.  
☒ Disabled  
☐ Use AWS Key Management Service (AWS KMS)  
☐ Use a hardware security module (HSM)

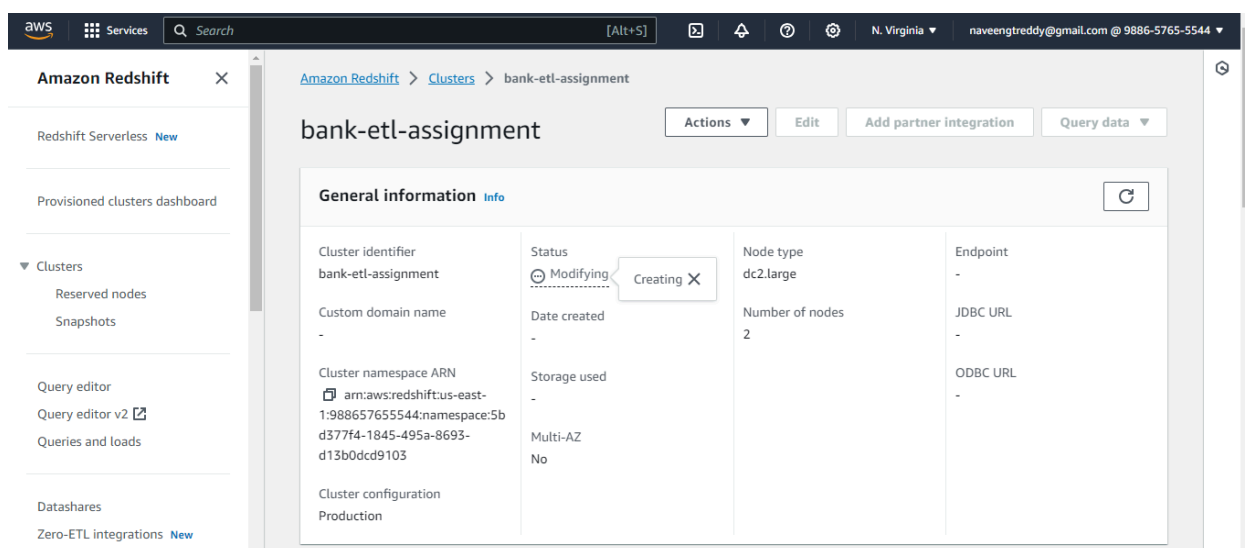
► **Maintenance** Info

► **Monitoring**

► **Backup**

Cancel **Create cluster**

## Redshift cluster under creation status



Amazon Redshift > Clusters > bank-etl-assignment

**bank-etl-assignment** Actions Edit Add partner integration Query data

**General information** Info

Cluster identifier bank-etl-assignment	Status Modifying Creating X	Node type dc2.large	Endpoint -
Custom domain name -	Date created -	Number of nodes 2	JDBC URL -
Cluster namespace ARN arn:aws:redshift:us-east-1:988657655544:namespace:5bd377f4-1845-495a-8693-d13b0dcd9103	Storage used -		ODBC URL -
Cluster configuration Production	Multi-AZ No		

## Redshift cluster under available status -

aws

Services

Search

[Alt+S]

N. Virginia

naveengtreddy@gmail.com @ 9886-5765-5544

Amazon Redshift

Redshift Serverless New

Provisioned clusters dashboard

Clusters

- Reserved nodes
- Snapshots

Query editor

Query editor v2

Queries and loads

Datashares

Zero-ETL integrations New

Amazon Redshift > Clusters > bank-etl-assignment

bank-etl-assignment

Actions Edit Add partner integration Query data

General information

Cluster identifier bank-etl-assignment	Status <span>Available</span>	Node type dc2.large	Endpoint bank-etl-assignment.cq6jeujwkqli.us-east-1.redshift.amazonaws.com:5439/redshiftdb
Custom domain name -	Date created February 04, 2024, 16:07 (UTC+05:30)	Number of nodes 2	JDBC URL jdbc:redshift://bank-etl-assignment.cq6jeujwkqli.us-east-1.redshift.amazonaws.com:5439/redshiftdb
Cluster namespace ARN arn:aws:redshift:us-east-1:988657655544:namespace:5bd377f4-1845-495a-8693-d13b0dcd9103	Storage used -		ODBC URL Driver={Amazon Redshift (x64)}; Server=bank-etl-assignment.cq6jeujwkqli.us-east-1.redshift.amazonaws.com; Database=redshiftrdb
Cluster configuration Production	Multi-AZ No		

# Setting up a database in the Redshift cluster and running queries to create the dimension and fact tables

## Loading data into a Redshift cluster from Amazon S3 bucket

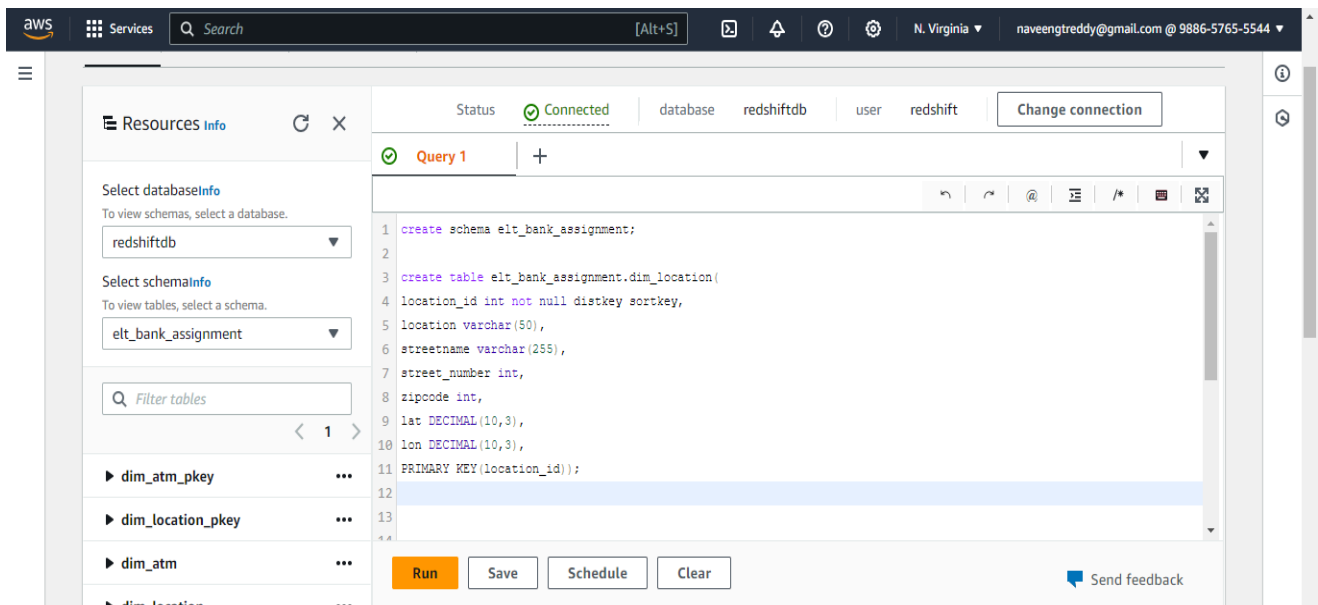
### Queries to create the dimension and fact tables with appropriate primary and foreign keys

- **Creating Schema :**

```
create schema elt_bank_assignment;
```

- **Creating location dimension Table.**

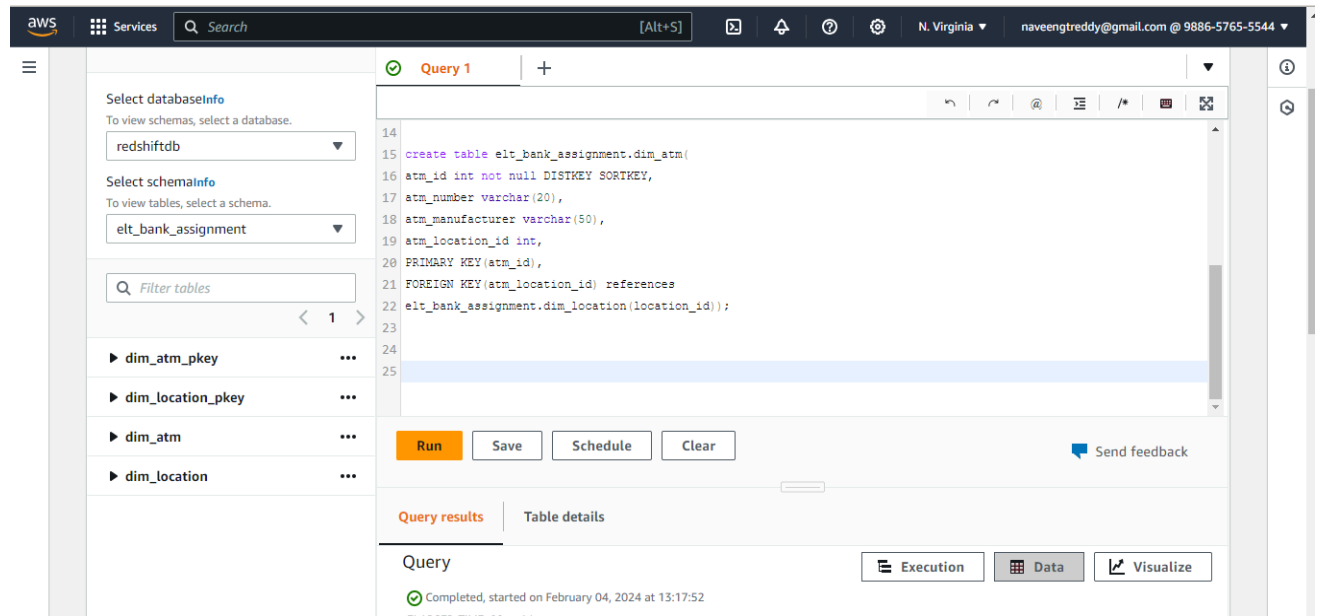
```
create table elt_bank_assignment.dim_location(
location_id int not null distkey sortkey,
location varchar(50),
streetname varchar(255),
street_number int,
zipcode int,
lat DECIMAL(10, 3),
lon DECIMAL(10, 3),
PRIMARY KEY(location_id)
);
```



The screenshot shows the AWS Redshift console interface. On the left, the 'Resources' sidebar is visible, showing the database 'redshiftdb' and the schema 'elt\_bank\_assignment'. The main area displays a SQL query being executed, which creates a schema 'elt\_bank\_assignment' and a table 'dim\_location' with columns: location\_id (int, not null, distkey sortkey), location (varchar(50)), streetname (varchar(255)), street\_number (int), zipcode (int), lat (DECIMAL(10,3)), lon (DECIMAL(10,3)), and a primary key on location\_id. The query is numbered 1 through 14. At the bottom, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The status bar at the top indicates 'Connected' to the 'redshiftdb' database using the 'redshift' user.

- **Creating ATM Dimention Table**

```
create table elt_bank_assignment.dim_atm(
atm_id int not null DISTKEY SORTKEY,
atm_number varchar(20),
atm_manufacturer varchar(50),
atm_location_id int,
PRIMARY KEY(atm_id),
FOREIGN KEY(atm_location_id) references
elt_bank_assignment.dim_location(location_id));
```

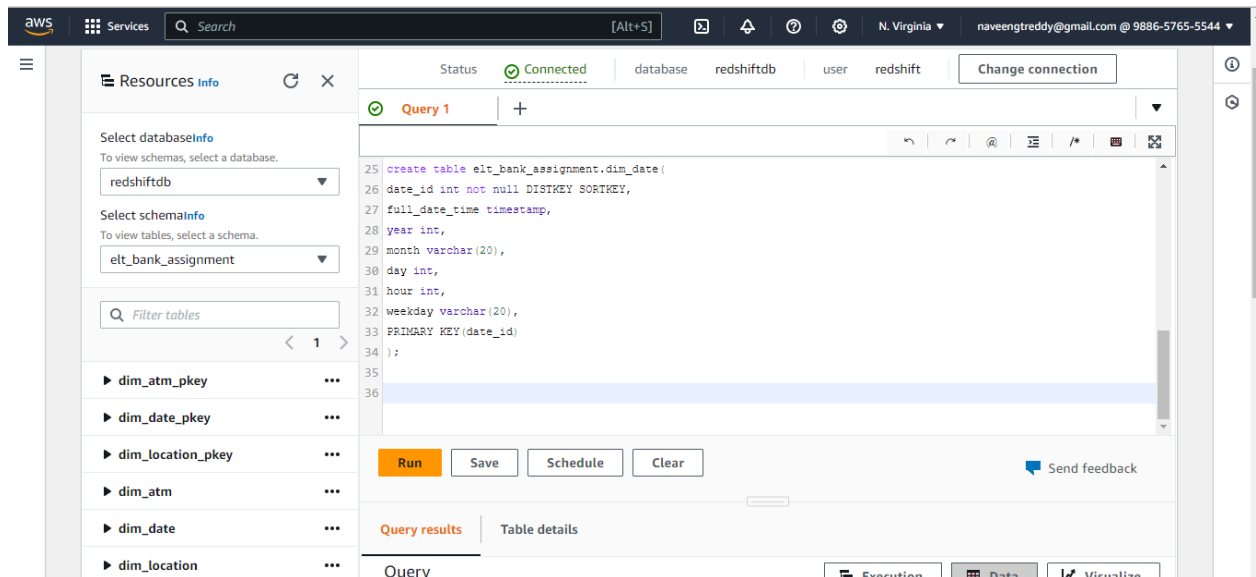


The screenshot shows the AWS Redshift console interface. On the left, the 'Select database' dropdown is set to 'redshiftdb' and the 'Select schema' dropdown is set to 'elt\_bank\_assignment'. Below these, a list of tables is shown: 'dim\_atm\_pkey', 'dim\_location\_pkey', 'dim\_atm', and 'dim\_location'. The main panel displays a SQL query in a text editor, which is the same query provided in the previous block. Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted. Below the buttons, the 'Query results' tab is selected, showing a message: 'Query Completed, started on February 04, 2024 at 13:17:52'. The 'Table details' tab is also visible.

- **Creating date dimension table**

```
create table elt_bank_assignment.dim_date(
date_id int not null DISTKEY SORTKEY,
full_date_time timestamp,
year int,
month varchar(20),
day int,
hour int,
weekday varchar(20),
PRIMARY KEY(date_id)
);
```





The screenshot shows the AWS Redshift console interface. On the left, the 'Resources Info' panel is open, showing the selected database as 'redshiftdb' and the schema as 'elt\_bank\_assignment'. Below this, a list of tables is displayed, including 'dim\_atm\_pkey', 'dim\_date\_pkey', 'dim\_location\_pkey', 'dim\_atm', 'dim\_date', and 'dim\_location'. The main query editor on the right contains the following SQL code:

```

25 create table elt_bank_assignment.dim_date (
26 date_id int not null DISTKEY SORTKEY,
27 full_date_time timestamp,
28 year int,
29 month varchar(20),
30 day int,
31 hour int,
32 weekday varchar(20),
33 PRIMARY KEY(date_id)
34 );

```

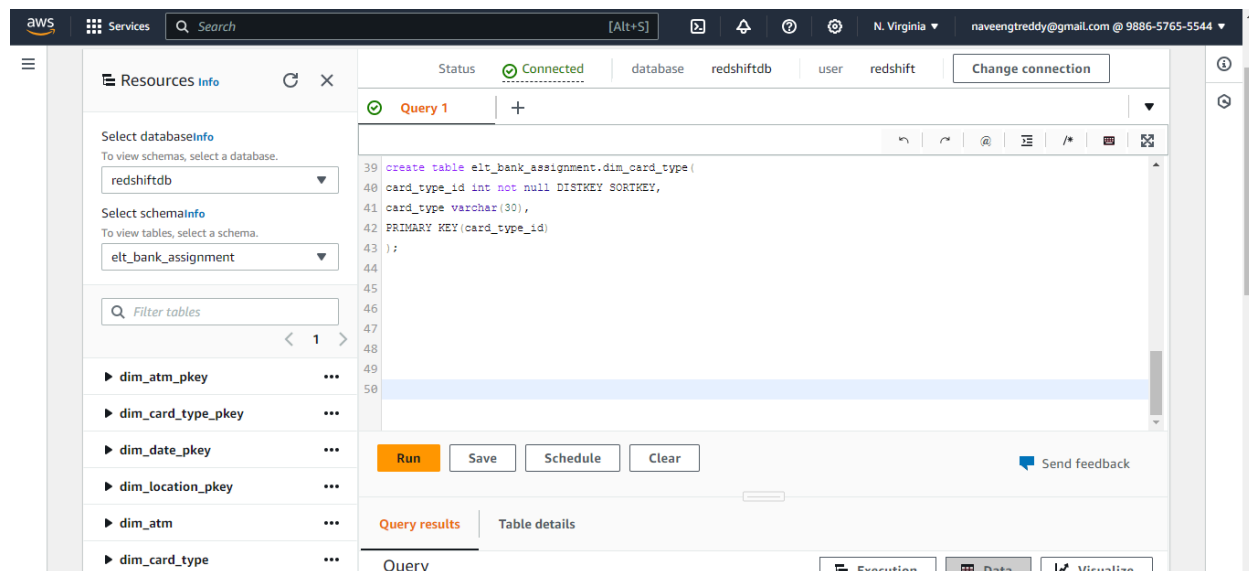
Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted in orange. At the bottom of the console, there are tabs for 'Query', 'Execution', 'Data', and 'Visualize'.

- **Creating card type dimension table**

```

create table elt_bank_assignment.dim_card_type (
card_type_id int not null DISTKEY SORTKEY,
card_type varchar(30),
PRIMARY KEY(card_type_id)
);

```



The screenshot shows the AWS Redshift console interface. On the left, the 'Resources Info' panel is open, showing the selected database as 'redshiftdb' and the schema as 'elt\_bank\_assignment'. Below this, a list of tables is displayed, including 'dim\_atm\_pkey', 'dim\_card\_type\_pkey', 'dim\_date\_pkey', 'dim\_location\_pkey', 'dim\_atm', and 'dim\_card\_type'. The main query editor on the right contains the following SQL code:

```

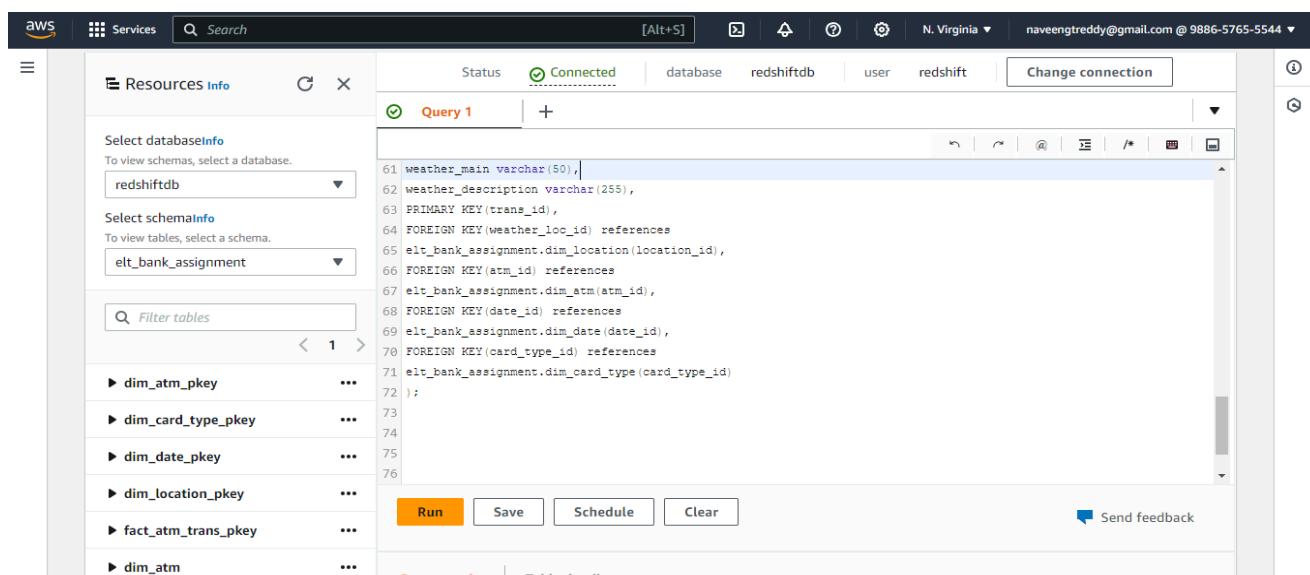
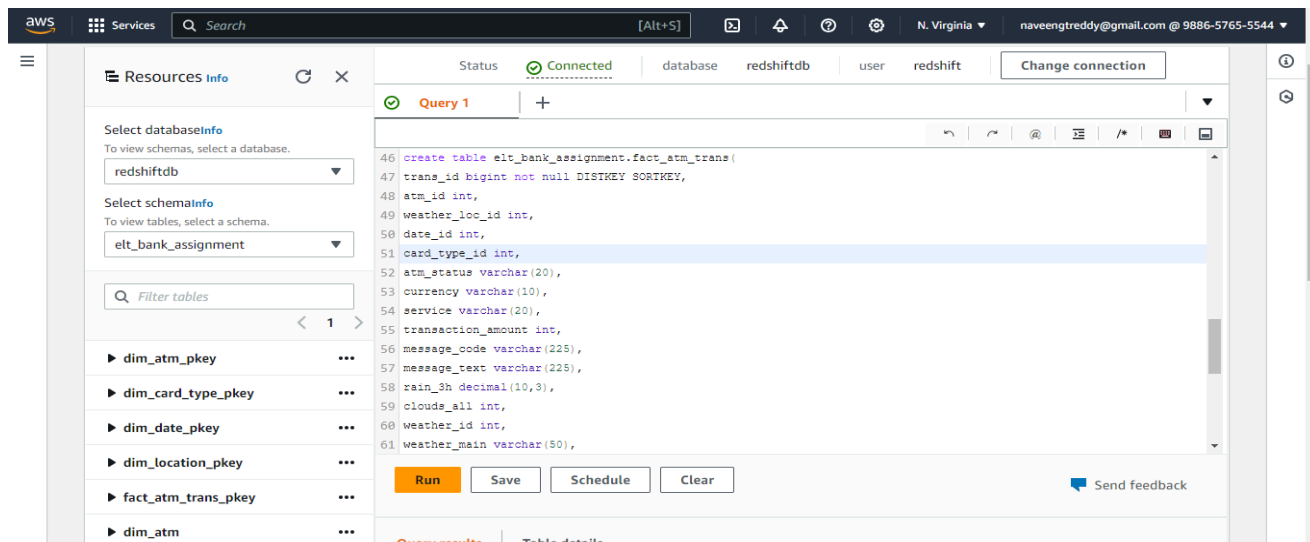
39 create table elt_bank_assignment.dim_card_type (
40 card_type_id int not null DISTKEY SORTKEY,
41 card_type varchar(30),
42 PRIMARY KEY(card_type_id)
43 );

```

Below the query editor, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Run' button is highlighted in orange. At the bottom of the console, there are tabs for 'Query', 'Execution', 'Data', and 'Visualize'.

- **Creating atm transactions fact table**

```
create table elt_bank_assignment.fact_atm_trans(  
  trans_id bigint not null DISTKEY SORTKEY,  
  atm_id int,  
  weather_loc_id int,  
  date_id int,  
  card_type_id int,  
  atm_status varchar(20),  
  currency varchar(10),  
  service varchar(20),  
  transaction_amount int,  
  message_code varchar(225),  
  message_text varchar(225),  
  rain_3h decimal(10,3),  
  clouds_all int,  
  weather_id int,  
  weather_main varchar(50),  
  weather_description varchar(255),  
  PRIMARY KEY(trans_id),  
  FOREIGN KEY(weather_loc_id) references  
  elt_bank_assignment.dim_location(location_id),  
  FOREIGN KEY(atm_id) references  
  elt_bank_assignment.dim_atm(atm_id),  
  FOREIGN KEY(date_id) references  
  elt_bank_assignment.dim_date(date_id),  
  FOREIGN KEY(card_type_id) references  
  elt_bank_assignment.dim_card_type(card_type_id)  
);
```

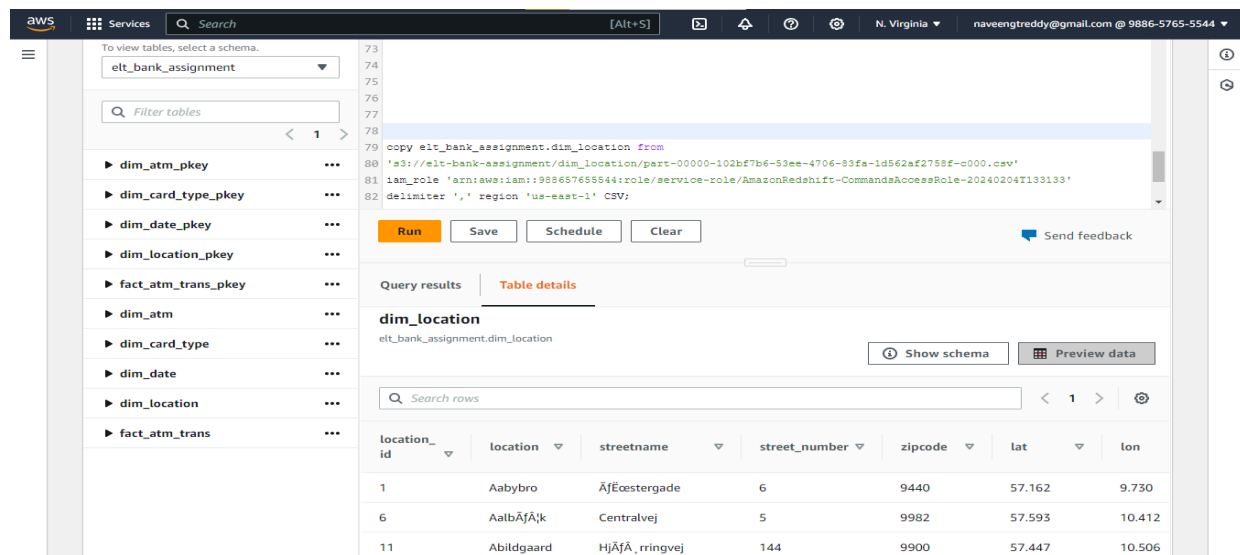


## Loading data into a Redshift cluster from Amazon S3 bucket

Queries to copy the data from S3 buckets to the Redshift cluster in the appropriate tables

- **Copying the data to dim\_location table**  
copy elt\_bank\_assignment.dim\_location from  
's3://elt-bank-assignment/dim\_location/part-00000-102bf7b6-53ee-4706-83fa-1d562af2758f-c000.csv'

```
iam_role
'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-Com
mandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1' CSV;
```

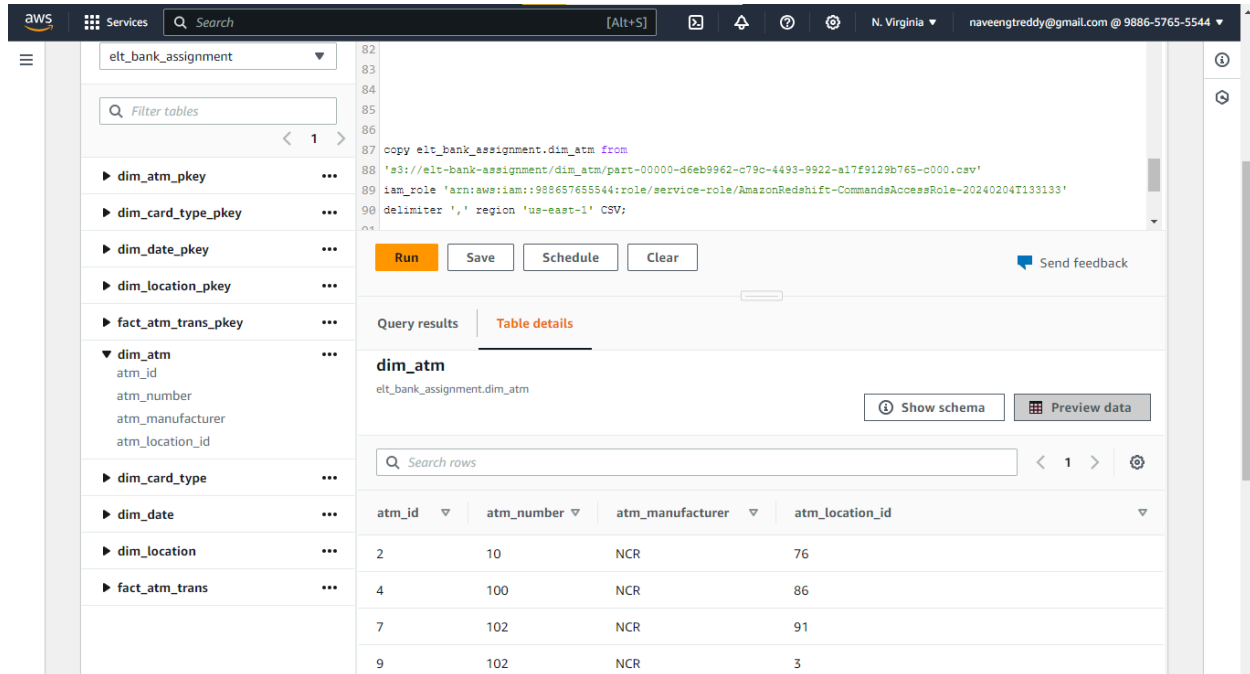


The screenshot shows the AWS Redshift console interface. On the left, a list of tables is visible, including 'dim\_atm\_pkey', 'dim\_card\_type\_pkey', 'dim\_date\_pkey', 'dim\_location\_pkey', 'fact\_atm\_trans\_pkey', 'dim\_atm', 'dim\_card\_type', 'dim\_date', 'dim\_location', and 'fact\_atm\_trans'. The main area displays a SQL query being executed: 'copy elt\_bank\_assignment.dim\_location from 's3://elt-bank-assignment/dim\_location/part-00000-102bf7b6-53ee-4706-83fa-1d562af2758f-0000.csv' iam\_role 'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-CommandsAccessRole-20240204T133133' delimiter ',' region 'us-east-1' CSV;'. Below the query, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. The 'Query results' tab is selected, showing the 'dim\_location' table with columns: location\_id, location, streetname, street\_number, zipcode, lat, and lon. The results show three rows of data.

location_id	location	streetname	street_number	zipcode	lat	lon
1	Aabybro	Åfjæstergade	6	9440	57.162	9.730
6	AalbÅfjÅk	Centralvej	5	9982	57.593	10.412
11	Abildgaard	HjÅfÅ_ringvej	144	9900	57.447	10.506

- **Copying the data to dim\_atm table**

```
copy elt_bank_assignment.dim_atm from
's3://elt-bank-assignment/dim_atm/part-00000-d6eb9962-c79c-4493
-9922-a17f9129b765-c000.csv'
iam_role
'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-Com
mandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1' CSV;
```



The screenshot shows the AWS Redshift console interface. On the left, a list of tables is displayed, including `dim_atm_pkey`, `dim_card_type_pkey`, `dim_date_pkey`, `dim_location_pkey`, `fact_atm_trans_pkey`, `dim_atm`, `dim_card_type`, `dim_date`, `dim_location`, and `fact_atm_trans`. The `dim_atm` table is selected, and its schema is shown: `atm_id`, `atm_number`, `atm_manufacturer`, and `atm_location_id`.

The query editor shows the following SQL query:

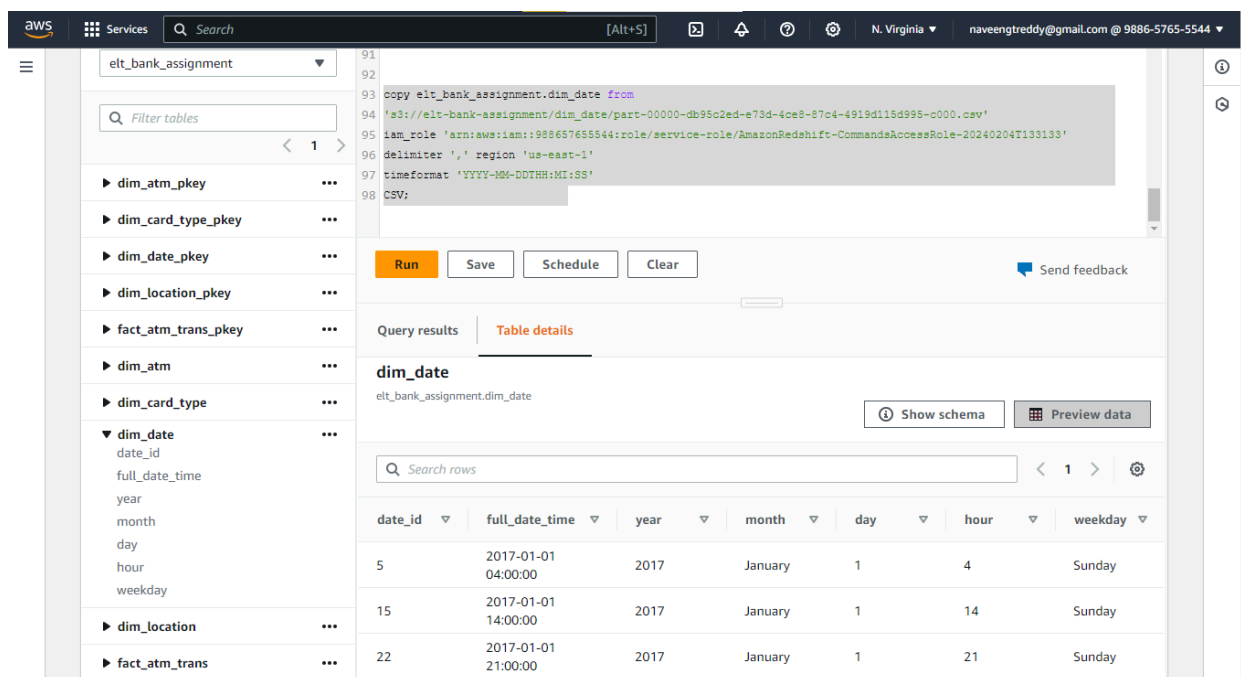
```
copy elt_bank_assignment.dim_atm from
's3://elt-bank-assignment/dim_atm/part-00000-d6eb9962-c79c-4493-9922-a17f9129b765-c000.csv'
iam_role 'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-CommandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1' CSV;
```

The query has been executed, and the results are displayed in the table below:

atm_id	atm_number	atm_manufacturer	atm_location_id
2	10	NCR	76
4	100	NCR	86
7	102	NCR	91
9	102	NCR	3

- **Copying the data to dim\_date table**

```
copy elt_bank_assignment.dim_date from
's3://elt-bank-assignment/dim_date/part-00000-db95c2ed-e73d-4ce
8-87c4-4919d115d995-c000.csv'
iam_role
'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-Com
mandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1'
timeformat 'YYYY-MM-DDTHH:MI:SS'
CSV;
```

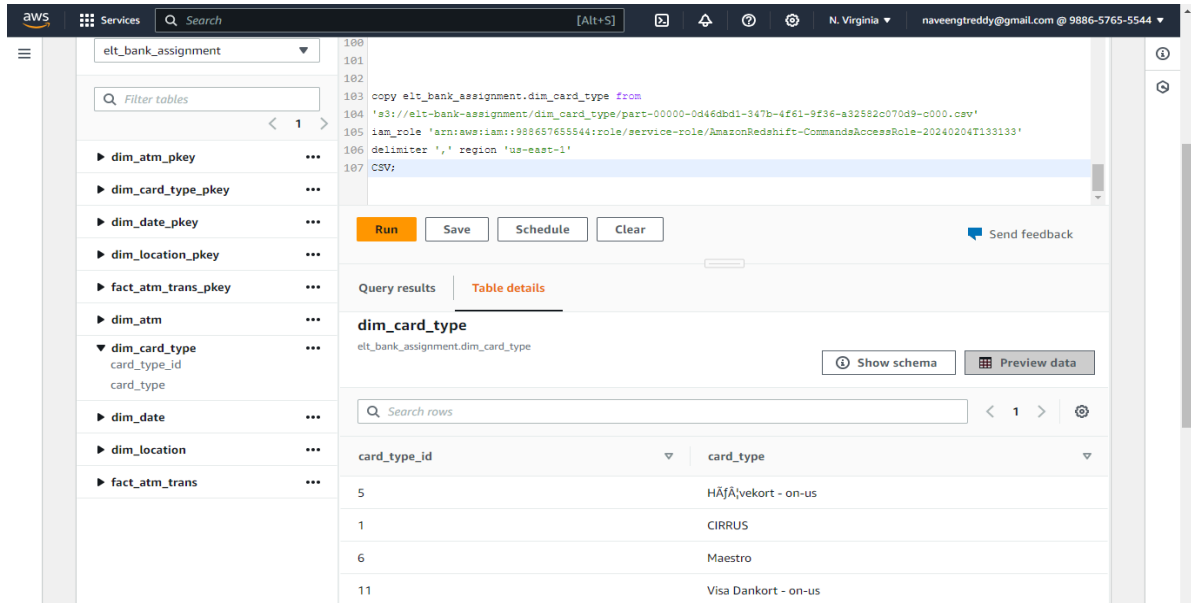


The screenshot shows the AWS Redshift console interface. On the left, a list of tables is displayed, including `dim_atm_pkey`, `dim_card_type_pkey`, `dim_date_pkey`, `dim_location_pkey`, `fact_atm_trans_pkey`, `dim_atm`, `dim_card_type`, `dim_date`, `dim_location`, and `fact_atm_trans`. The `dim_date` table is selected, and its schema is shown: `date_id`, `full_date_time`, `year`, `month`, `day`, `hour`, and `weekday`. The query execution results are displayed in a table format:

date_id	full_date_time	year	month	day	hour	weekday
5	2017-01-01 04:00:00	2017	January	1	4	Sunday
15	2017-01-01 14:00:00	2017	January	1	14	Sunday
22	2017-01-01 21:00:00	2017	January	1	21	Sunday

## ● Copying the data to dim\_card\_type table

```
copy elt_bank_assignment.dim_card_type from
's3://elt-bank-assignment/dim_card_type/part-00000-0d46dbd1-347
b-4f61-9f36-a32582c070d9-c000.csv'
iam_role
'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-Com
mandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1'
CSV;
```



The screenshot shows the AWS Redshift console interface. On the left, a list of tables is displayed, including dim\_atm\_pkey, dim\_card\_type\_pkey, dim\_date\_pkey, dim\_location\_pkey, fact\_atm\_trans\_pkey, dim\_atm, dim\_card\_type, dim\_date, dim\_location, and fact\_atm\_trans. The 'dim\_card\_type' table is selected. The main panel shows the query execution results for the 'dim\_card\_type' table. The query is: `copy elt_bank_assignment.dim_card_type from 's3://elt-bank-assignment/dim_card_type/part-00000-0d46dbd1-347b-4f61-9f36-a32582c070d9-c000.csv' iam_role 'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-CommandsAccessRole-20240204T133133' delimiter ',' region 'us-east-1' CSV;`. The results are displayed in a table with columns 'card\_type\_id' and 'card\_type'. The data rows are: (5, HÅfÅ\vekort - on-us), (1, CIRRUS), (6, Maestro), and (11, Visa Dankort - on-us).

- **Copying the data to fact\_atm\_trans table**

```
copy elt_bank_assignment.fact_atm_trans from
's3://elt-bank-assignment/fact_atm_trans/part-00000-4d319008-f8
f8-4bdf-a63a-32287e3ecf39-c000.csv'
iam_role
'arn:aws:iam::988657655544:role/service-role/AmazonRedshift-Com
mandsAccessRole-20240204T133133'
delimiter ',' region 'us-east-1'
CSV;
```

© Copyright. upGrad Education Pvt. Ltd. All rights reserved